

```
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import pandas as pd
```

```
df = sns.load_dataset("titanic")
df.head()
```

	survived	pclass	sex	age	sibsp	parch	fare	embarked	class	who
0	0	3	male	22.0	1	0	7.2500	S	Third	man
1	1	1	female	38.0	1	0	71.2833	C	First	woman
2	1	3	female	26.0	0	0	7.9250	S	Third	woman
3	1	1	female	35.0	1	0	53.1000	S	First	woman
4	0	3	male	35.0	0	0	8.0500	S	Third	man

Next steps:

[Generate code with df](#)

[New interactive sheet](#)

```
df = df[['age', 'fare', 'survived']]
```

```
df.head()
```

	age	fare	survived	
0	22.0	7.2500	0	
1	38.0	71.2833	1	
2	26.0	7.9250	1	
3	35.0	53.1000	1	
4	35.0	8.0500	0	

Next steps:

[Generate code with df](#)

[New interactive sheet](#)

```

from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score
from sklearn.preprocessing import KBinsDiscretizer
from sklearn.model_selection import cross_val_score
from sklearn.compose import ColumnTransformer

```

```

df.dropna(inplace=True)
df.isnull().sum()

```

```

      0
age    0
fare    0
survived 0

```

dtype: int64

```
df.head()
```

	age	fare	survived	
0	22.0	7.2500	0	
1	38.0	71.2833	1	
2	26.0	7.9250	1	
3	35.0	53.1000	1	
4	35.0	8.0500	0	

Next steps:

[Generate code with df](#)
[New interactive sheet](#)

```

X = df.iloc[:,0:2]
Y = df.iloc[:, -1]

```

```
X.head()
```

	age	fare
0	22.0	7.2500
1	38.0	71.2833
2	26.0	7.9250
3	35.0	53.1000
4	35.0	8.0500

Next steps:

[Generate code with X](#)[New interactive sheet](#)

```
Y.head()
```

	survived
0	0
1	1
2	1
3	1
4	0

dtype: int64

```
X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size=0.4,random_state
```

```
clf = DecisionTreeClassifier()  
clf.fit(X_train,Y_train)
```

▼ DecisionTreeClassifier ⓘ ?

```
DecisionTreeClassifier()
```

```
y_pred = clf.predict(X_test)
```

```
accuracy_score(Y_test,y_pred)
```

```
0.5524475524475524
```

```
np.mean(cross_val_score(DecisionTreeClassifier(),X,Y,cv=5,scoring="accuracy"))
```

```
np.float64(0.6176499556781246)
```

```
kb_age = KBinsDiscretizer(n_bins=15,encode="ordinal",strategy="quantile")
kb_fire = KBinsDiscretizer(n_bins=10,encode="ordinal",strategy="quantile")
```

```
trf = ColumnTransformer(
    [
        ("first",kb_age,[0]),
        ("second",kb_fire,[1])
    ]
)
```

```
X_train_trf = trf.fit_transform(X_train)
X_test_trf = trf.fit_transform(X_test)
```

```
trf.named_transformers_["first"].n_bins_
array([15])
```

```
trf.named_transformers_["second"].n_bins_
array([10])
```

```
trf.named_transformers_["first"].bin_edges_
array([array([ 0.42,  6.   , 16.   , 18.   , 20.   , 22.   , 24.   , 26.   , 28.   ,
              30.   , 32.   , 35.   , 38.   , 44.   , 52.   , 71.   ]),
      dtype=object)
```

```
trf.named_transformers_["second"].bin_edges_
array([array([ 0.      ,  7.69585,  7.925  ,  9.28335, 12.2875 , 13.92915,
              21.      , 27.7354 , 35.5    , 72.25   , 512.3292 ]),
      dtype=object)
```

```
df1 = pd.DataFrame({
    "age" : X_train["age"],
    "age_trf" : X_train_trf[:,0],
    "fare" : X_train["fare"],
    "fare_trf" : X_train_trf[:,1]
})
```

```
df1.head()
```

	age	age_trf	fare	fare_trf	
146	27.0	6.0	7.7958	1.0	
155	51.0	13.0	61.3792	8.0	
735	28.5	7.0	16.1000	5.0	
294	24.0	5.0	7.8958	2.0	
0	22.0	4.0	7.2500	0.0	

Next steps:

[Generate code with df1](#)[New interactive sheet](#)

```
df1["age_label"] = pd.cut(x=X_train["age"],bins=trf.named_transformers_["first"]
df1["fare_label"] = pd.cut(x=X_train["fare"],bins=trf.named_transformers_["secc
```

df1.head()

	age	age_trf	fare	fare_trf	age_label	fare_label	
146	27.0	6.0	7.7958	1.0	(26.0, 28.0]	(7.696, 7.925]	
155	51.0	13.0	61.3792	8.0	(44.0, 52.0]	(35.5, 72.25]	
735	28.5	7.0	16.1000	5.0	(28.0, 30.0]	(13.929, 21.0]	
294	24.0	5.0	7.8958	2.0	(22.0, 24.0]	(7.696, 7.925]	
0	22.0	4.0	7.2500	0.0	(20.0, 22.0]	(0.0, 7.696]	

Next steps:

[Generate code with df1](#)[New interactive sheet](#)

```
clf = DecisionTreeClassifier()
clf.fit(X_train_trf,Y_train)
```

▼ DecisionTreeClassifier ⓘ ?

```
DecisionTreeClassifier()
```

```
y_pred2 = clf.predict(X_test_trf)
y_pred2
```

```
array([0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 1,
       0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0,
       0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 1, 1, 0, 0, 1,
       1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0,
       0, 1, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1, 1, 0, 0,
       1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 1, 0, 1, 1, 0,
       0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
       0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0,
       0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 1, 0, 0, 1, 1, 1, 0, 0, 1, 0, 0,
       1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0,
       0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1,
       1, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0,
       0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0])
```

```
accuracy_score(Y_test,y_pred2)
```

```
0.6293706293706294
```

```
trf_test = trf.fit_transform(X)
```

```
np.mean(cross_val_score(DecisionTreeClassifier(),X,Y,cv=5,scoring="accuracy"))
```

```
np.float64(0.6190485570767261)
```

```
def discretized(bins, strategy, X_train, X_test, Y_train):
    # Create discretizers
    kb_age = KBinsDiscretizer(n_bins=bins, encode="ordinal", strategy=strategy)
    kb_fire = KBinsDiscretizer(n_bins=bins, encode="ordinal", strategy=strategy)

    # ColumnTransformer
    trf = ColumnTransformer([
        ("age_bin", kb_age, [0]),
        ("fire_bin", kb_fire, [1])
    ])

    # Transform training and testing sets
    X_train_trf = trf.fit_transform(X_train)
    X_test_trf = trf.transform(X_test)

    # Train classifier
    clf = DecisionTreeClassifier()
    clf.fit(X_train_trf, Y_train)

    # Predict
    y_pred2 = clf.predict(X_test_trf)
    test_acc = accuracy_score(Y_test, y_pred2)
    print("Test Accuracy:", test_acc)

    # Cross-validation on full transformed data
    X_trf_full = trf.fit_transform(np.vstack((X_train, X_test)).astype(float))
```

```
Y_full = np.hstack((Y_train, Y_test))
cv_acc = np.mean(cross_val_score(DecisionTreeClas:
print("Cross-validation Accuracy:", cv_acc)
```

```
plt.figure(figsize=(14,4))
plt.subplot(121)
plt.hist(X["age"])
plt.title("Age Before Discretization")
```

```
plt.subplot(122)
plt.hist( X_trf_full[:,0],color = "red")
plt.title("Age after Discretization")
```

```
plt.figure(figsize=(14,4))
plt.subplot(121)
plt.hist(X["fare"])
plt.title("fare Before Discretization")
```


```
plt.subplot(122)
plt.hist( X_trf_full[:,1],color = "g")
plt.title("fare after Discretization")
```

```
return clf, test_acc, cv_acc
```

```
clf, test_acc, cv_acc = discretized(10, "quantile", X_train, X_test, Y_train, Y
```

Test Accuracy: 0.6433566433566433

Cross-validation Accuracy: 0.67374175120654



```
def discretized(bins, strategy, X_train, X_test, Y_train, Y_test):
    # Create discretizers
    kb_age = KBinsDiscretizer(n_bins=bins, encode="ordinal", strategy=strategy)
```

```

kb_fire = KBinsDiscretizer(n_bins=bins, encode="ordinal", strategy=strategy)

# ColumnTransformer
trf = ColumnTransformer([
    ("age_bin", kb_age, [0]),
    ("fire_bin", kb_fire, [1])
])

# Transform training and testing sets
X_train_trf = trf.fit_transform(X_train)
X_test_trf = trf.transform(X_test)

# Train classifier
clf = DecisionTreeClassifier()
clf.fit(X_train_trf, Y_train)

# Predict
y_pred2 = clf.predict(X_test_trf)
test_acc = accuracy_score(Y_test, y_pred2)
print("Test Accuracy:", test_acc)

# Cross-validation on full transformed data
X_trf_full = trf.fit_transform(np.vstack((X_train, X_test)))
Y_full = np.hstack((Y_train, Y_test))
cv_acc = np.mean(cross_val_score(DecisionTreeClassifier(), X_trf_full, Y_full))
print("Cross-validation Accuracy:", cv_acc)

plt.figure(figsize=(14,4))
plt.subplot(121)
plt.hist(X["age"])
plt.title("Age Before Discretization")

plt.subplot(122)
plt.hist( X_trf_full[:,0],color = "red")
plt.title("Age after Discretization")

plt.figure(figsize=(14,4))
plt.subplot(121)
plt.hist(X["fare"])
plt.title("fare Before Discretization")


plt.subplot(122)
plt.hist( X_trf_full[:,1],color = "g")
plt.title("fare after Discretization")

return clf, test_acc, cv_acc

```



```
clf, test_acc, cv_acc = discretized(5, "uniform", X_train, X_test, Y_train, Y_test)
Test Accuracy: 0.6538461538461539
```



```
clf, test_acc, cv_acc = discretized(10, "kmeans", X_train, X_test, Y_train, Y_test)
```

Test Accuracy: 0.6328671328671329

Cross-validation Accuracy: 0.6499162809021963

