

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler,MinMaxScaler,RobustScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC

```

```
df = pd.read_csv("/content/Social_Network_Ads (1).csv")
```

```
df.head()
```

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0

```
round(df.describe(),1)
```

	User ID	Age	EstimatedSalary	Purchased
count	400.0	400.0	400.0	400.0
mean	15691539.8	37.7	69742.5	0.4
std	71658.3	10.5	34097.0	0.5
min	15566689.0	18.0	15000.0	0.0
25%	15626763.8	29.8	43000.0	0.0
50%	15694341.5	37.0	70000.0	0.0
75%	15750363.0	46.0	88000.0	1.0
max	15815236.0	60.0	150000.0	1.0

```
df = df.iloc[:, 2:]
```

```
df.head()
```

	Age	EstimatedSalary	Purchased
0	19	19000	0
1	35	20000	0
2	26	43000	0
3	27	57000	0
4	19	76000	0

```
X_train,X_test,y_train,y_test = train_test_split(df.drop("Purchased",axis =1),
                                                df["Purchased"],test_size=0.3,random_state=0)
```

```
X_train.shape,X_test.shape
```

```
((280, 2), (120, 2))
```

```
X_train.head()
```

	Age	EstimatedSalary
92	26	15000
223	60	102000
234	38	112000
232	40	107000
377	42	53000

```
X_test.head()
```

	Age	EstimatedSalary
132	30	87000
309	38	50000
341	35	75000
196	30	79000
246	35	50000

```
Scaler = StandardScaler()
Scaler.fit(X_train)
```

▼ StandardScaler ⓘ ?

```
StandardScaler()
```

```
X_train_scaled = Scaler.transform(X_train)
X_test_scaled = Scaler.transform(X_test)
```

```
X_train_scaled = pd.DataFrame(X_train_scaled,columns = X_train.columns)
X_test_scaled = pd.DataFrame(X_test_scaled, columns = X_test.columns)
```

```
X_train_scaled.head()
```

	Age	EstimatedSalary
0	-1.163172	-1.584970
1	2.170181	0.930987
2	0.013305	1.220177
3	0.209385	1.075582
4	0.405465	-0.486047

```
X_test_scaled.head()
```

	Age	EstimatedSalary
0	-0.771013	0.497201
1	0.013305	-0.572804
2	-0.280814	0.150172
3	-0.771013	0.265849
4	-0.280814	-0.572804

```
round(X_train_scaled.describe(),1)
```


	Age	EstimatedSalary
count	280.0	280.0
mean	0.0	0.0
std	1.0	1.0
min	-1.9	-1.6
25%	-0.8	-0.8
50%	-0.1	0.0
75%	0.8	0.5
max	2.2	2.3

```
round(X_test_scaled.describe(),1)
```


	Age	EstimatedSalary
count	120.0	120.0
mean	-0.1	-0.0
std	1.1	1.0
min	-1.9	-1.6
25%	-1.0	-0.7
50%	-0.2	-0.2
75%	0.7	0.5
max	2.2	2.3

```
# Effect of Scaling
```

```
fig,(ax1,ax2) = plt.subplots(ncols = 2,figsize = (12,5))
ax1.scatter(X_train["Age"],X_train["EstimatedSalary"],color = "red")
plt.xlabel("Age")
plt.ylabel("Salary")
ax1.set_title("before Scalling")
ax2.scatter(X_train_scaled["Age"],X_train_scaled["EstimatedSalary"],color = "g")
plt.xlabel("Age")
plt.ylabel("Salary")
ax2.set_title("after Scalling")
plt.show()
```



```
fig,(ax1,ax2) = plt.subplots(ncols = 2,figsize = (12,5))
ax1.scatter(X_test["Age"],X_test["EstimatedSalary"],color = "red")
plt.xlabel("Age")
plt.ylabel("Salary")
ax1.set_title("before Scalling")
ax2.scatter(X_test_scaled["Age"],X_test_scaled["EstimatedSalary"],color = "g")
plt.xlabel("Age")
plt.ylabel("Salary")
ax2.set_title("after Scalling")
plt.show()
```




```
fig, (ax1, ax2) = plt.subplots(ncols=2, figsize=(12,5))

# Before Scaling
ax1.set_title("Before Scaling")
sns.kdeplot(X_train["Age"], ax=ax1, label="Age")
sns.kdeplot(X_train["EstimatedSalary"], ax=ax1, label="Estimated Salary")
ax1.legend()

# After Scaling
ax2.set_title("After Scaling")
sns.kdeplot(X_train_scaled["Age"], ax=ax2, label="Age (Scaled)")
sns.kdeplot(X_train_scaled["EstimatedSalary"], ax=ax2, label="Estimated Salary (Scaled)")
ax2.legend()

plt.show()
```



```
fig, (ax1, ax2) = plt.subplots(ncols=2, figsize=(12,5))
```


```

# Before Scaling
ax1.set_title("Before Scaling")
sns.kdeplot(X_test["Age"], ax=ax1, label="Age")
sns.kdeplot(X_test["EstimatedSalary"], ax=ax1, label="Estimated Salary")
ax1.legend()

# After Scaling
ax2.set_title("After Scaling")
sns.kdeplot(X_test_scaled["Age"], ax=ax2, label="Age (Scaled)")
sns.kdeplot(X_test_scaled["EstimatedSalary"], ax=ax2, label="Estimated Salary (Scaled)")
ax2.legend()

plt.show()


```




```

# comprision of distribution
fig, (ax1, ax2) = plt.subplots(ncols=2, figsize=(12,5))
# Before Scaling
ax1.set_title("Before Scaling")
sns.kdeplot(X_train["Age"],ax = ax1,color="r")
# After Scaling
ax2.set_title("After Scaling")
sns.kdeplot(X_train_scaled["Age"],ax = ax2,color="g")
plt.show()

```



```
# comprision of distribution
fig, (ax1, ax2) = plt.subplots(ncols=2, figsize=(12,5))
# Before Scaling
ax1.set_title("Before Scaling")
sns.kdeplot(X_train["EstimatedSalary"],ax = ax1,color="r")
# After Scaling
ax2.set_title("After Scaling")
sns.kdeplot(X_train_scaled["EstimatedSalary"],ax = ax2,color="g")
plt.show()
```



```
# Logistic_model

log = LogisticRegression()
log_scaled = LogisticRegression()
```

```
log.fit(X_train,y_train)
log_scaled.fit(X_train_scaled,y_train)
```

▼ LogisticRegression ⓘ ?
 LogisticRegression()

```
y_pred = log.predict(X_test)
Y_pred_scaled = log_scaled.predict(X_test_scaled)
```

```
print("Accuracy_score before scaling",accuracy_score(y_test, y_pred))
print("Accuracy_score After Scaling",accuracy_score(y_test, Y_pred_scaled))
```

```
Accuracy_score before scaling 0.875
Accuracy_score After Scaling 0.8666666666666667
```

DECISION TREE

```
dt = DecisionTreeClassifier()
dt_Scaled = DecisionTreeClassifier()
```

```
dt.fit(X_train,y_train)
dt_Scaled.fit(X_train_scaled,y_train)
```

▼ DecisionTreeClassifier ⓘ ?
 DecisionTreeClassifier()

```
y_preddt = dt.predict(X_test)
y_preddt_scaled = dt_Scaled.predict(X_test_scaled)
```

```
print("Actual accuracy score before scaling",accuracy_score(y_test,y_preddt))
print("scaled accuracy score after scaling ",accuracy_score(y_test,y_preddt_scaled))
```

```
Actual accuracy score before scaling 0.875
scaled accuracy score after scaling 0.875
```

```
model = RandomForestClassifier(n_estimators=100, random_state=42)
model_scaled = RandomForestClassifier(n_estimators=100, random_state=42)
```

```
model.fit(X_train,y_train)
model_scaled.fit(X_train_scaled,y_train)
```

▼ RandomForestClassifier ⓘ ?
 RandomForestClassifier(random_state=42)

```
y_predRF = model.predict(X_test)
y_predRF_scaled = model_scaled.predict(X_test_scaled )
```

```
print("Actual Accoracy before scaling",accuracy_score(y_test,y_predRF))
print("scaled Accuracy after scaling ",accuracy_score(y_test,y_predRF_scaled))
```

```
Actual Accoracy before scaling 0.9166666666666666
scaled Accuracy after scaling 0.9166666666666666
```

```
knnmodel = KNeighborsClassifier(n_neighbors=5)
knnmodel_scaled = KNeighborsClassifier(n_neighbors=5)
```

```
knnmodel.fit(X_train,y_train)
knnmodel_scaled.fit(X_train_scaled,y_train)
```

▼ KNeighborsClassifier ⓘ ?

```
KNeighborsClassifier()
```

```
y_predknn = knnmodel.predict(X_test)
y_predknn_scaled = knnmodel_scaled.predict(X_test_scaled)
```

```
print("Actual Accoracy before scaling",accuracy_score(y_test,y_predknn))
print("scaled Accuracy after scaling ",accuracy_score(y_test,y_predknn_scaled))
```

```
Actual Accoracy before scaling 0.825
scaled Accuracy after scaling 0.9166666666666666
```

```
SVM1 = SVC()
SVM1_scaled = SVC()
```

```
SVM1.fit(X_train,y_train)
SVM1_scaled.fit(X_train_scaled,y_train)
```

▼ SVC ⓘ ?

```
SVC()
```

```
y_predSVM = SVM1.predict(X_test)
y_predSVM_scaled = SVM1_scaled.predict(X_test_scaled)
```

```
print("Actual Accoracy before scaling",accuracy_score(y_test,y_predSVM))
print("scaled Accuracy after scaling ",accuracy_score(y_test,y_predSVM_scaled))
```

```
Actual Accoracy before scaling 0.7916666666666666
scaled Accuracy after scaling 0.9083333333333333
```

```
# With Outlier
```


```
df = pd.concat(
    [
        df,
        pd.DataFrame({
            "Age": [5, 80, 90],
            "EstimatedSalary": [1000, 30000, 70000],
            "Purchased": [0, 1, 1]
        })
    ],
    ignore_index=True
)
```

```
df.head()
```

	Age	EstimatedSalary	Purchased
0	19	19000	0
1	35	20000	0
2	26	43000	0
3	27	57000	0
4	19	76000	0

```
plt.scatter(df["Age"],df["EstimatedSalary"])
```

<matplotlib.collections.PathCollection at 0x7adb18d82ba0>



```
X_train,X_test,y_train,y_test = train_test_split(df.drop("Purchased",axis =1),
                                                  df["Purchased"],test_size=0.3,random_state=0)
```

```
X_train.shape,X_test.shape
```

```
((282, 2), (121, 2))
```

```
Scaler = StandardScaler()
Scaler.fit(X_train)
```

▼ StandardScaler ⓘ ?


StandardScaler()

```
X_train_scaled = Scaler.transform(X_train)
X_test_scaled = Scaler.transform(X_test)
```


```
X_train_scaled = pd.DataFrame(X_train_scaled,columns = X_train.columns)
X_test_scaled = pd.DataFrame(X_test_scaled, columns = X_test.columns)
```

```
# Effect of Scaling
```

```
fig,(ax1,ax2) = plt.subplots(ncols = 2,figsize = (12,5))
ax1.scatter(X_train["Age"],X_train["EstimatedSalary"],color = "red")
plt.xlabel("Age")
plt.ylabel("Salary")
ax1.set_title("before Scalling")
ax2.scatter(X_train_scaled["Age"],X_train_scaled["EstimatedSalary"],color = "g")
plt.xlabel("Age")
plt.ylabel("Salary")
ax2.set_title("after Scalling")
plt.show()
```



```
fig,(ax1,ax2) = plt.subplots(ncols = 2,figsize = (12,5))
ax1.scatter(X_test["Age"],X_test["EstimatedSalary"],color = "red")
plt.xlabel("Age")
plt.ylabel("Salary")
ax1.set_title("before Scalling")
ax2.scatter(X_test_scaled["Age"],X_test_scaled["EstimatedSalary"],color = "g")
plt.xlabel("Age")
plt.ylabel("Salary")
ax2.set_title("after Scalling")
plt.show()
```




```
fig, (ax1, ax2) = plt.subplots(ncols=2, figsize=(12,5))

# Before Scaling
ax1.set_title("Before Scaling")
sns.kdeplot(X_train["Age"], ax=ax1, label="Age")
sns.kdeplot(X_train["EstimatedSalary"], ax=ax1, label="Estimated Salary")
ax1.legend()

# After Scaling
ax2.set_title("After Scaling")
sns.kdeplot(X_train_scaled["Age"], ax=ax2, label="Age (Scaled)")
sns.kdeplot(X_train_scaled["EstimatedSalary"], ax=ax2, label="Estimated Salary (Scaled)")
ax2.legend()

plt.show()
```



```
fig, (ax1, ax2) = plt.subplots(ncols=2, figsize=(12,5))

# Before Scaling
```


```

ax1.set_title("Before Scaling")
sns.kdeplot(X_test["Age"], ax=ax1, label="Age")
sns.kdeplot(X_test["EstimatedSalary"], ax=ax1, label="Estimated Salary")
ax1.legend()

# After Scaling
ax2.set_title("After Scaling")
sns.kdeplot(X_test_scaled["Age"], ax=ax2, label="Age (Scaled)")
sns.kdeplot(X_test_scaled["EstimatedSalary"], ax=ax2, label="Estimated Salary (Scaled)")
ax2.legend()

plt.show()

```



```
# Logistic_model
```

```
log = LogisticRegression()
log_scaled = LogisticRegression()
```

```
log.fit(X_train,y_train)
log_scaled.fit(X_train_scaled,y_train)
```

▼ LogisticRegression ⓘ ?

```
LogisticRegression()
```

```
y_pred = log.predict(X_test)
Y_pred_scaled = log_scaled.predict(X_test_scaled)
```

```
print("Accuracy_score before scaling",accuracy_score(y_test, y_pred))
print("Accuracy_score After Scaling",accuracy_score(y_test, Y_pred_scaled))
```

```
Accuracy_score before scaling 0.859504132231405
Accuracy_score After Scaling 0.8512396694214877
```

```
dt = DecisionTreeClassifier()
dt_Scaled = DecisionTreeClassifier()
```

```
dt.fit(X_train,y_train)
dt_Scaled.fit(X_train_scaled,y_train)
```

▼ DecisionTreeClassifier ⓘ ?

```
DecisionTreeClassifier()
```

```
y_preddt = dt.predict(X_test)
y_preddt_scaled = dt_Scaled.predict(X_test_scaled)
```

```
print("Actual accuracy score before scaling",accuracy_score(y_test,y_preddt))
print("scaled accuracy score after scaling ",accuracy_score(y_test,y_preddt_scaled))
```

```
Actual accuracy score before scaling 0.8760330578512396
scaled accuracy score after scaling 0.8760330578512396
```

```
model = RandomForestClassifier(n_estimators=100, random_state=42)
model_scaled = RandomForestClassifier(n_estimators=100, random_state=42)
```

```
model.fit(X_train,y_train)
model_scaled.fit(X_train_scaled,y_train)
```

▼ RandomForestClassifier ⓘ ?

```
RandomForestClassifier(random_state=42)
```

```
y_predRF = model.predict(X_test)
y_predRF_scaled = model_scaled.predict(X_test_scaled )
```

```
print("Actual accuracy score before scaling",accuracy_score(y_test,y_predRF))
print("scaled accuracy score after scaling ",accuracy_score(y_test,y_predRF_scaled))
```

```
Actual accuracy score before scaling 0.8925619834710744
scaled accuracy score after scaling 0.8925619834710744
```

```
knnmodel = KNeighborsClassifier(n_neighbors=5)
knnmodel_scaled = KNeighborsClassifier(n_neighbors=5)
```

```
knnmodel.fit(X_train,y_train)
knnmodel_scaled.fit(X_train_scaled,y_train)
```

▼ KNeighborsClassifier ⓘ ?

```
KNeighborsClassifier()
```

```
y_predknn = knnmodel.predict(X_test)
y_predknn_scaled = knnmodel_scaled.predict(X_test_scaled)
```

```
print("Actual Accoracy before scaling",accuracy_score(y_test,y_predknn))
print("scaled Accuracy after scaling ",accuracy_score(y_test,y_predknn_scaled))
```